

Would you trust this preview?

It looks like a real photo.
But the shirt has changed color.

How closely can AI reproduce
the clothes we choose?



Requested garment




Pretrained preview

Selected example: 07703_00, seed 9026. One example, not a failure-rate estimate.

Better garment reconstruction, without unwanted changes elsewhere?



Edited person input



Requested garment



Original reference


1,000 training cases

64 validation cases

32 final-test cases

VITON-HD + VITON-HD-edit. Custom reconstruction study for clothing-preview developers.

I adapted the existing attention weights



1,000 cases, three epochs, 3,000 optimizer updates

AdamW, learning rate 0.00001. Training objective: predicted-noise MSE.

Both models get the same test



Garment MAE and SSIM

How close is the garment to the reference?

Outside-clothing MAE

What changed outside the clothing?


32 cases × 2 models × 2 seeds

128 attempts, 122 images, 27 complete paired cases.

The safety checker excluded six attempts across five cases. No retries.

384 × 512 pixels, 20 DDIM steps, guidance 2.5. Average seeds within each case.

Garment error fell 15.6% on average



Each line is one case. Lower error is better.

21 / 27

cases improved

Mean MAE

0.1335 pretrained


0.1127 fine-tuned

Green lines improve.

Six red lines worsen.

Paired mean difference: -0.02085 . Descriptive 95% case-bootstrap interval: $[-0.03222, -0.01032]$.

Preservation did not improve for every person



14 / 27

cases had more
outside-clothing error

Yet average outside error
fell 25.9%.

Larger gains in some cases
outweighed smaller regressions.

Each dot is one paired case. Differences are fine-tuned minus pretrained. Lower left improves both errors.

Even the largest gain still changes the person



Input person



Pretrained



Fine-tuned



Reference

The clothing is closer to the reference.

The arm pose and trousers still change.

Case 08931_00, seed 9026. Selected by the largest case-average garment-MAE reduction. Full gallery includes all cases.

A closer clothing match. Still unwanted changes elsewhere.

Attention fine-tuning helped on this dataset.
It did not make every clothing preview dependable.

Klara shows where fine-tuning helps
and where the model still makes mistakes.

Limits: edited inputs, 27 complete pairs, one training seed.
Identity separation and upstream exposure are unverified.

klara-app.de

Live demo uses the trained checkpoint.
Recorded gallery works offline.