

Would you trust this preview?

It looks like a real photo.
But the shirt has changed color.

How closely can AI reproduce
the clothes we choose?



Requested garment



Pretrained preview

Selected example: 07703_00, seed 9026. One example, not a failure-rate estimate.

Can attention fine-tuning improve the garment while preserving the person and background?



Edited person input



Requested garment



Original reference

VITON-HD and VITON-HD-edit provide an aligned reference for evaluation.
For developers building clothing previews for online shoppers.

Separate data for learning and evaluation



Checks before modeling

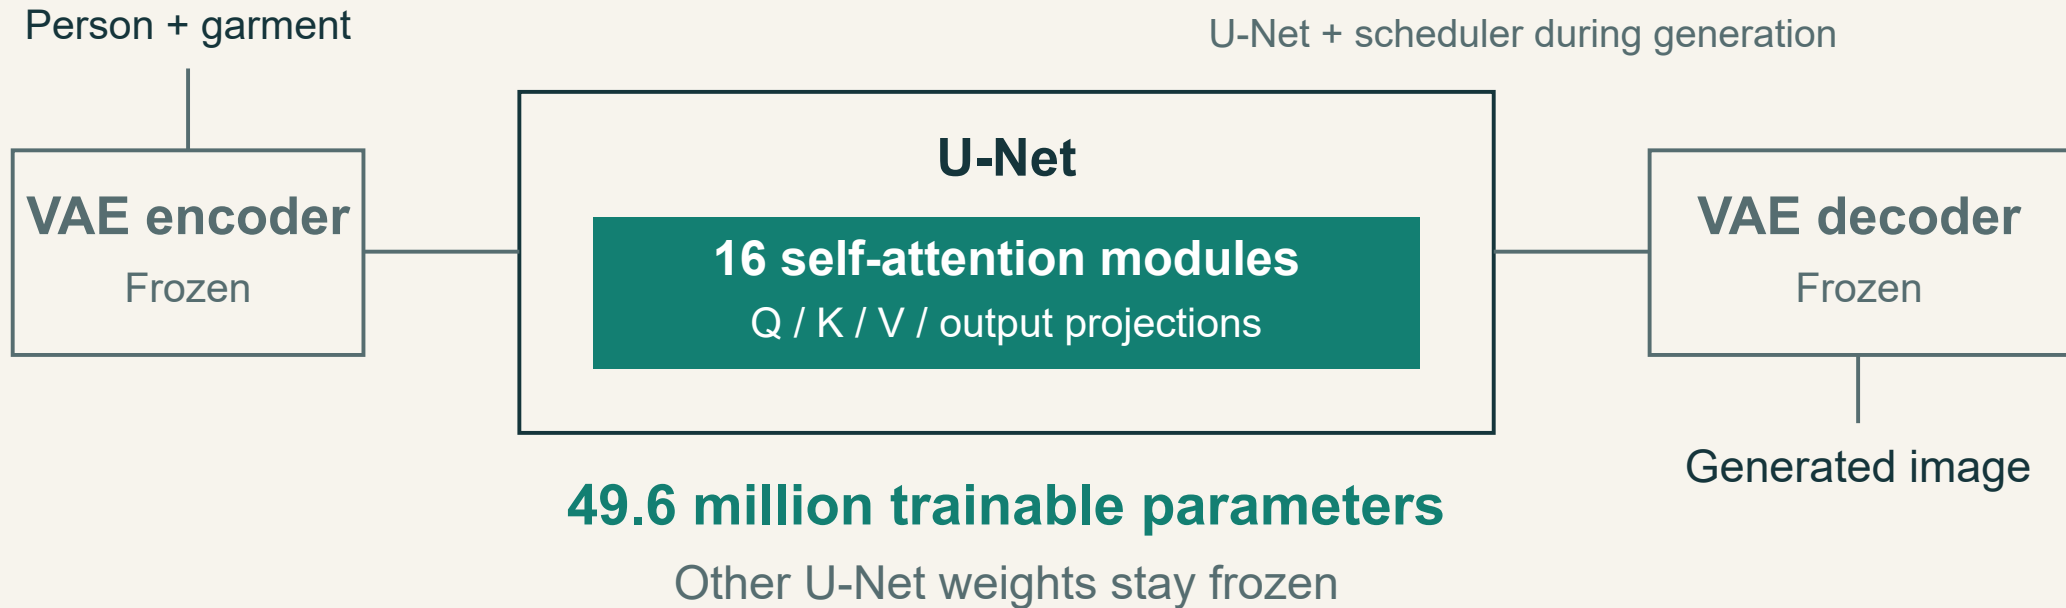
Pairings and image dimensions
Missing files and exact duplicates
Source IDs across the splits

What these checks cannot tell me

Different source IDs do not prove
that every person is different.

Custom split of the original VITON-HD test partition.
Upstream training exposure is unverified.

Fine-tuning the attention already in the model



The experiment changes selected attention weights.
The VAE and other U-Net weights stay frozen.

A convincing image can still be wrong



Inside the clothing region

MAE measures pixel differences.

SSIM measures structural agreement.

Outside the clothing region

Measure unwanted changes
relative to the input person.

Same cases, seeds and settings for both models.

Masks define measurement regions, not generator inputs.

The feedback I need

Do these measurements answer the research question clearly?

What should I explain more simply for a mixed technical and nontechnical jury?

After the MVP: review the method and findings, then prepare the final story and demonstration.

Rehearsal: 14 October. Jury: 15 October.